

## Online Learning

In batch learning task, a model is learned from a batch of data. In online learning task, a model is continuously learned (or, updated) from a sequence of data. Let  $x_1, x_2, \dots, x_n$  be a sequence of instances, and  $y_t$  be the label of  $x_t$ . A general framework of online learning is as follows:

- [0] initialize  $f$
- [1]  $f$  receives  $x_t$  at time  $t$
- [2]  $f$  predicts label  $\hat{y}_t$  for  $x_t$
- [3]  $f$  receives true label  $y_t$
- [4]  $f$  updates itself based on some loss  $\ell(y_t, \hat{y}_t)$
- [5] repeat step 1-4 until convergence

After  $n$  rounds,  $f$  will incur a cumulative loss

$$L_n(f) = \frac{1}{n} \sum_{t=1}^n \ell(\hat{y}_t, y_t). \quad (1)$$

The goal of online learning is to find a model  $f$  that can minimize  $L_n(f)$ . The challenge is one does not store all received data and continuously use them to retrain the model – that's too inefficient in terms of computation and memory. Instead, one can keep updating the model using newly received instances – this can be implemented using stochastic gradient descent (SGD).

Consider the following loss

$$L_n(f) = \sum_{t=1}^n \ell(y_t, f(x_t)). \quad (2)$$

Ordinary gradient descent updates  $f$  using the gradient of  $L_n$  on all instances, i.e.,

$$f = f - \eta \cdot \frac{\partial}{\partial f} L_n(f) = f - \eta \cdot \sum_{t=1}^n \frac{\partial}{\partial f} \ell(y_t, f(x_t)). \quad (3)$$

SGD approximates (3) using the gradient on a single instance  $x_t$  at time  $t$ , i.e.,

$$f = f - \eta \cdot \sum_{t=1}^n \frac{\partial}{\partial f} \ell(y_t, f(x_t)) \approx f - \eta \cdot \frac{\partial}{\partial f} \ell(y_t, f(x_t)). \quad (4)$$

Therefore, Step 4 in the online learning framework can be implemented as

- [4]  $f$  updates itself based on  $f = f - \eta \cdot \frac{\partial}{\partial f} \ell(\hat{y}_t; y_t)$ .

Stochastic gradient descent will converge if  $f$  is a bit more than convex. In general, it converges more slowly than ordinary gradient descent.

## The Peceptron Algorithm

Consider binary classification task with label set  $Y = \{-1, +1\}$ . Let  $\text{sign}(z)$  be a sign function outputting  $+1$  if  $z \geq 0$  and  $-1$  otherwise. Perceptron is an online learning algorithm for model

$$f(x) = \text{sign}(x^T \beta). \quad (5)$$

It finds an  $f$  that can minimize the following hinge loss (that only counts misclassified instances)

$$L_n(f) = \sum_{t=1}^n \max\{0, -y_t(x_t^T \beta)\}. \quad (6)$$

[*Discussion*] How does the hinge loss count only misclassified instances?

Let  $I_{mis}$  be the index set of misclassified instances. The gradient of  $L_n$  is

$$L'_n(f) = \sum_{t \in I_{mis}} -y_t x_t. \quad (7)$$

[*Exercise*] Derive (7).

Perceptron applies SGD to optimize  $L_n$  and therefore approximates (7) by

$$L'_n \approx -y_t x_t, \quad \text{if } x_t \text{ is misclassified.} \quad (8)$$

Perceptron learning is summarized in Algorithm 1. A common choice of learning rate is  $\eta = 1$ .

---

### Algorithm 1 The Perceptron Learning Algorithm

---

**Input:** learning rate  $\eta$

**Initialization:** (randomly) initialize model  $\beta$

**for**  $t = 1$  to  $n$  **do**

- 1: model receives instance  $x_t$
- 2: model predicts  $\hat{y}_t = \text{sign}(x_t^T \beta)$
- 3: model receives true label  $y_t$
- 4: if  $\hat{y}_t \neq y_t$ , then model updates itself by

$$\beta = \beta - \eta \cdot (-y_t \cdot x_t) = \beta + \eta y_t x_t. \quad (9)$$

otherwise, model does nothing.

**end for**

---